

# Hadoop Cluster Commands

[https://www.youtube.com/watch?v=\\_zR0v4XIFNs](https://www.youtube.com/watch?v=_zR0v4XIFNs)

```
## on all nodes  
sudo apt update
```

```
sudo apt install ssh  
sudo apt install pdsh
```

```
nano .bashrc  
## add to the end of .bashrc file of all nodes  
export PDSH_RCMD_TYPE=ssh
```

```
## create keygen on all three nodes and skip any input req using enter  
ssh-keygen -t rsa -P ""
```

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

```
## login to localhost and enter yes in req  
ssh localhost  
sudo apt install openjdk-8-jdk
```

```
sudo apt install openjdk-11-jdk
```

```
java -version
```

```
sudo wget -P ~  
https://mirrors.sonic.net/apache/hadoop/common/hadoop-3.2.1/hadoop-3.2.1.tar.gz
```

## Hadoop 3.3.1

```
Alternative : sudo wget -P ~  
https://mirror.olnevhost.net/pub/apache/hadoop/common/hadoop-3.3.1/hadoop-3.3.1.tar.gz
```

```
tar xzf hadoop-3.3.1.tar.gz  
mv hadoop-3.3.1 hadoop
```

```
tar xzf hadoop-3.2.1.tar.gz  
mv hadoop-3.2.1 hadoop  
nano ~/hadoop/etc/hadoop/hadoop-env.sh
```

Goes at the top of Hadoop-env.sh file

```
export JAVA_HOME=/usr/lib/jvm/java-11-openjdk-amd64/
```

```
sudo mv hadoop /usr/local/hadoop
```

```
sudo nano /etc/environment
```

# put the following path below the existing path in the /etc/environment

```
PATH="/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games:/usr/local/games:/usr/local/hadoop/bin:/usr/local/hadoop/sbin"JAVA_HOME="/usr/lib/jvm/java-11-openjdk-amd64/jre"
```

## enter 5 commands below on all the nodes and enter password twice and press enter for

## everyother option

```
[hadoop-master@hadoop-master:~$ tar xzf hadoop-3.3.1.tar.gz
[hadoop-master@hadoop-master:~$ mv hadoop-3.3.1 hadoop
[hadoop-master@hadoop-master:~$ nano ~/hadoop/etc/hadoop/hadoop-env.sh
[hadoop-master@hadoop-master:~$ sudo mv hadoop /usr/local/hadoop
[hadoop-master@hadoop-master:~$ sudo nano /etc/environment
hadoop-master@hadoop-master:~$ sudo adduser hadoopuser
Adding user `hadoopuser' ...
Adding new group `hadoopuser' (1001) ...
Adding new user `hadoopuser' (1001) with group `hadoopuser' ...
Creating home directory `/home/hadoopuser' ...
Copying files from `/etc/skel' ...
[Enter new UNIX password:
[Retype new UNIX password:
passwd: password updated successfully
Changing the user information for hadoopuser
Enter the new value, or press ENTER for the default
[         Full Name []:
[         Room Number []:
[         Work Phone []:
[         Home Phone []:
[         Other []:
[Is the information correct? [Y/n]
hadoop-master@hadoop-master:~$
```

```
sudo adduser hadoopuser
```

```
sudo usermod -aG hadoopuser hadoopuser
```

```
sudo chown hadoopuser:root -R /usr/local/hadoop/
```

```
sudo chmod g+rwX -R /usr/local/hadoop/
```

```
sudo adduser hadoopuser sudo
```

# to find the ip addresses of each node then paste that info in /etc/hosts  
ip addr

```
10.0.0.4 hadoop-master
10.0.0.5 hadoop-slave1
10.0.0.6 hadoop-slave2
```

sudo nano /etc/hosts

```
10.0.0.4 hadoop-master
10.0.0.5 hadoop-slave1
10.0.0.6 hadoop-slave2
```

sudo nano /etc/hostname  
sudo reboot

# login to hadoopuser on master-node only  
su - hadoopuser  
# for master-node only and press enter for  
ssh-keygen -t rsa

# enter following one by one only on master and enter yes and enter password  
ssh-copy-id hadoopuser@hadoop-master  
ssh-copy-id hadoopuser@hadoop-slave1  
ssh-copy-id hadoopuser@hadoop-slave2

# run below command block individually on slave 1 and slave 2  
sudo adduser hadoopuser  
sudo usermod -aG hadoopuser hadoopuser  
sudo chown hadoopuser:root -R /usr/local/hadoop/  
sudo chmod g+rxw -R /usr/local/hadoop/  
sudo adduser hadoopuser sudo

sudo nano /usr/local/hadoop/etc/hadoop/core-site.xml

```
<property>
<name>fs.defaultFS</name>
```

```
<value>hdfs://hadoop-master:9000</value>
</property>
```

```
sudo nano /usr/local/hadoop/etc/hadoop/hdfs-site.xml
```

```
<property>
<name>dfs.namenode.name.dir</name><value>/usr/local/hadoop/data/nameNode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name><value>/usr/local/hadoop/data/dataNode</value>
</property>
<property>
<name>dfs.replication</name>
<value>2</value>
</property>
```

```
# only on hadoop-master
sudo nano /usr/local/hadoop/etc/hadoop/workers
```

```
hadoop-slave1
hadoop-slave2
```

```
scp /usr/local/hadoop/etc/hadoop/* hadoop-slave1:/usr/local/hadoop/etc/hadoop/
scp /usr/local/hadoop/etc/hadoop/* hadoop-slave2:/usr/local/hadoop/etc/hadoop/
```

```
## only on master node
source /etc/environment
cd /usr/local/hadoop/
bin/hdfs namenode -format
sbin/start-dfs.sh
```

```
## now go to azure cloud web portal and in the overview section of hadoop-master change dns
name to hadoop-master-1765
```

## go in networking of hadoop-master and add inbound port rule. Change port name and destination port ranges to 9870

## go in networking of hadoop-master and add inbound port rule. Change port name and destination port ranges to 8088

## go to hadoop-master copy dns name and paste it in the search bar  
Enter : url:9870

## only on master node

```
export HADOOP_HOME="/usr/local/hadoop"  
export HADOOP_COMMON_HOME=$HADOOP_HOME  
export HADOOP_CONF_DIR=$HADOOP_HOME/etc/hadoop  
export HADOOP_HDFS_HOME=$HADOOP_HOME  
export HADOOP_MAPRED_HOME=$HADOOP_HOME  
export HADOOP_YARN_HOME=$HADOOP_HOME
```

```
sudo nano /usr/local/hadoop/etc/hadoop/yarn-site.xml
```

```
<property>  
<name>yarn.resourcemanager.hostname</name>  
<value>hadoop-master</value>  
</property>
```

```
sbin/start-yarn.sh
```

```
sbin/start-all.sh
```

```
scp /usr/local/hadoop/etc/hadoop/* hadoop-slave1:/usr/local/hadoop/etc/hadoop/  
scp /usr/local/hadoop/etc/hadoop/* hadoop-slave2:/usr/local/hadoop/etc/hadoop/
```

```
sbin/stop-all.sh
```

```
sbin/start-all.sh
```

```
hadoop fs -put /Users/wajeeh/Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/data.txt /user
```

```
hadoop fs -put /Users/wajeeh/Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/data.txt /user
```

```
scp data.txt hadoopuser@hadoop-master:"/Users/wajeeh/Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/data.txt"
```

```
"/Users/wajeeh/Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/data.txt"
```

```
hadoop fs -mkdir /user
```

```
cat data.txt |python3 /Users/wajeeh/Desktop/Big Data Assignments/A3/Big-Data-Assignment-3/mapper.py |python3 /Users/wajeeh/Desktop/Big Data Assignments/A3/Big-Data-Assignment-3/reducer.py^C
```

```
hadoop fs -put data.txt /user
```

```
hadoop jar /home/talal/hadoop-3.3.4/share/hadoop/tools/lib/hadoop-streaming-3.3.4.jar -file /home/talal/mapper1.py -mapper 'python3 mapper1.py' -file /home/talal/reducer1.py -reducer 'python3 reducer1.py' -input /user/file.txt -output /user/output4  
hadoop fs -rm -r /user/output5
```

```
# uploading mapper.py sample.txt and data.txt and reducer.py from local machine to  
hadoop-master at default /home/hadoop-master
```

```
scp Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/sample.txt Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/mapper.py Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/reducer.py hadoop-master@20.110.56.176:~
```

```
scp Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/sample.txt Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/mapper.py Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/reducer.py hadoop-slave1@20.110.153.135:~
```

```
scp Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/sample.txt Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/mapper.py Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/reducer.py hadoop-slave2@20.110.159.120:~
```

```
scp Desktop/Big\ Data\ Assignments/A3/Big-Data-Assignment-3/sample.txt
hadoop-master@20.110.56.176:~
```

```
# uploading mapper.py sample.txt and reducer.py from default /home/hadoop-master to
hadoop-master at default /home/hadoopuser to make it accessible
```

```
scp ~/mapper.py hadoopuser@hadoop-master:~
scp ~/reducer.py hadoopuser@hadoop-master:~
scp ~/sample.txt hadoopuser@hadoop-master:~
```

```
scp ~/mapper.py hadoopuser@hadoop-slave1:~
scp ~/reducer.py hadoopuser@hadoop-slave1:~
scp ~/sample.txt hadoopuser@hadoop-slave1:~
```

```
scp ~/mapper.py hadoopuser@hadoop-slave2:~
scp ~/reducer.py hadoopuser@hadoop-slave2:~
scp ~/sample.txt hadoopuser@hadoop-slave2:~
```

```
## /path/to/hadoop-streaming.jar
find / -name "hadoop-streaming*.jar" 2>/dev/null
```

In out case it was as follows

```
/usr/local/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.3.1.jar
```

```
# run mapreduce
```

```
hadoop jar /usr/local/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.3.1.jar -file mapper.py -mapper mapper.py
-file reducer.py -reducer reducer.py -input /user/sample.txt -output /user/result
```

```
# copy the output.txt created in hadoopuser to home directory of hadoop-master
sudo scp hadoopuser@hadoop-master:output.txt /home/hadoop-master
```

```
# copy the output.txt from home directory of hadoop-master to local machine
```

```
scp hadoop-master@20.110.56.176:output.txt Desktop/Big\ Data\  
Assignments/A3/Big-Data-Assignment-3
```